TA Session 2

2018.10.15

Wonjong Rhee, Hyunghun Cho, Daeyoung Choi
Seoul National University
Graduate School of Convergence Science and Technology
Applied Data Science Lab.

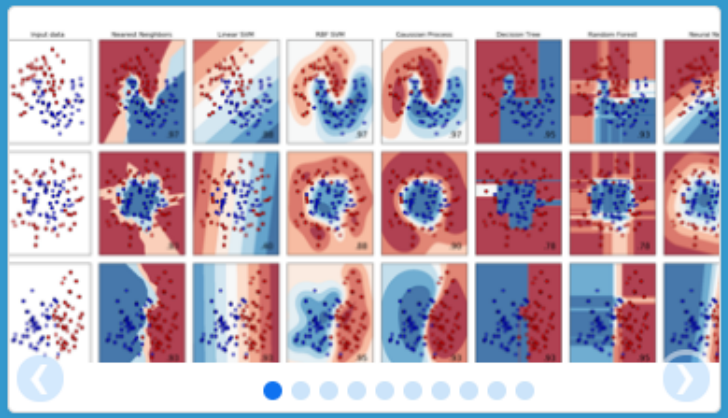

scikit-learn

Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which category an object belongs to.

Applications: Spam detection, Image recognition.
Algorithms: SVM, nearest neighbors, random forest, ... — Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.
Algorithms: SVR, ridge regression, Lasso, ... — Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes
Algorithms: k-Means, spectral clustering, mean-shift, ... — Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency
Algorithms: PCA, feature selection, non-negative matrix factorization. — Examples

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning
Modules: grid search, cross validation, metrics. — Examples

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.
Modules: preprocessing, feature extraction.

— Examples

scikit-learn

sklearn.model_selection: Model Selection

User guide: See the Cross-validation: evaluating estimator performance, Tuning the hyper-parameters of an estimator and Learning curve sections for further details.

Splitter Classes

model_selection.GroupKFold([n_splits])	K-fold iterator variant with non-overlapping groups.
model_selection.GroupShuffleSplit([...])	Shuffle-Group(s)-Out cross-validation iterator
model_selection.KFold([n_splits, shuffle, ...])	K-Folds cross-validator
model_selection.LeaveOneGroupOut()	Leave One Group Out cross-validator
model_selection.LeavePGroupsOut(n_groups)	Leave P Group(s) Out cross-validator
model_selection.LeaveOneOut()	Leave-One-Out cross-validator
model_selection.LeavePOut(p)	Leave-P-Out cross-validator
model_selection.PredefinedSplit(test_fold)	Predefined split cross-validator
model_selection.RepeatedKFold([n_splits, ...])	Repeated K-Fold cross validator.
model_selection.RepeatedStratifiedKFold([...])	Repeated Stratified K-Fold cross validator.
model_selection.ShuffleSplit([n_splits, ...])	Random permutation cross-validator
model_selection.StratifiedKFold([n_splits, ...])	Stratified K-Folds cross-validator
model_selection.StratifiedShuffleSplit([...])	Stratified ShuffleSplit cross-validator
model_selection.TimeSeriesSplit([n_splits, ...])	Time Series cross-validator

scikit-learn

sklearn.ensemble : Ensemble Methods

The `sklearn.ensemble` module includes ensemble-based methods for classification, regression and anomaly detection.

User guide: See the Ensemble methods section for further details.

`ensemble.AdaBoostClassifier` ([...])	An AdaBoost classifier.
`ensemble.AdaBoostRegressor` ([base_estimator, ...])	An AdaBoost regressor.
`ensemble.BaggingClassifier` ([base_estimator, ...])	A Bagging classifier.
`ensemble.BaggingRegressor` ([base_estimator, ...])	A Bagging regressor.
`ensemble.ExtraTreesClassifier` ([...])	An extra-trees classifier.
`ensemble.ExtraTreesRegressor` ([n_estimators, ...])	An extra-trees regressor.
`ensemble.GradientBoostingClassifier` ([loss, ...])	Gradient Boosting for classification.
`ensemble.GradientBoostingRegressor` ([loss, ...])	Gradient Boosting for regression.
`ensemble.IsolationForest` ([n_estimators, ...])	Isolation Forest Algorithm
`ensemble.RandomForestClassifier` ([...])	A random forest classifier.
`ensemble.RandomForestRegressor` ([...])	A random forest regressor.
`ensemble.RandomTreesEmbedding` ([...])	An ensemble of totally random trees.
`ensemble.VotingClassifier` (estimators[, ...])	Soft Voting/Majority Rule classifier for unfitted estimators.

scikit-learn

Model validation

model_selection.cross_validate (estimator, X)	Evaluate metric(s) by cross-validation and also record fit/score times.
model_selection.cross_val_predict (estimator, X)	Generate cross-validated estimates for each input data point
model_selection.cross_val_score (estimator, X)	Evaluate a score by cross-validation
model_selection.learning_curve (estimator, X, y)	Learning curve.
model_selection.permutation_test_score (...)	Evaluate the significance of a cross-validated score with permutations
model_selection.validation_curve (estimator, ...)	Validation curve.

Task 1: LOOCV and K-fold cross validation

- ISLR/Chapter 5.ipynb 을 실행하고 결과를 확인해 보세요.

Task 2: Tree based methods

- ISLR/Chapter 8.ipynb 을 실행하고 결과를 확인해 보세요.

Task 3: combining Task 1 and 2

- 1. Chapter 8.ipynb에 있는 Boston dataset을 이용하여 Chapter 5.ipynb의 Figure 5.2와 Figure 5.4 (오른쪽 10-fold CV만) 를 재현하세요. 알고리즘은 RandomForest를 사용하며 10-fold CV, hyperparameter는 feature의 개수(1~10개)로 합니다.

- 2. 1번에 해당하는 ISLR의 text를 읽고 1번의 결과를 설명하시오. 책의 관련 내용 (178p, 182p와 그 주변 page들) 을 읽고 이해하여 적으시면 됩니다.

- 3. Boston dataset으로 RandomForest 모델을 만든 후 다음과 같이 feature importance 그림을 그려보세요. Feature의 개수, tree의 개수, data split은 자유입니다.

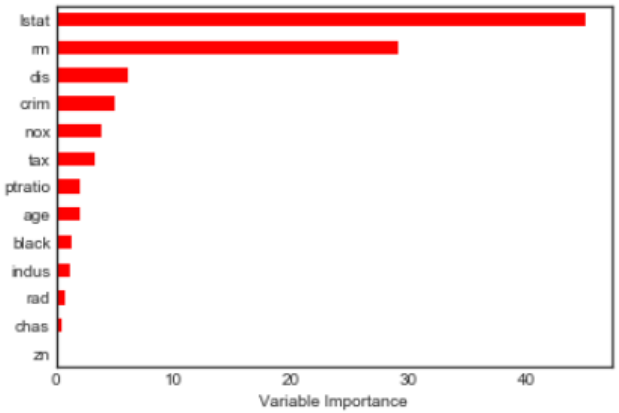

- 4. (advanced) 1번을 GradientBoostingRegressor functio을 이용해 재현해보세요. Hyperparameter는 n_estimators로 합니다. (range는 자유)

유의사항

- Task는 구현 자체보다는 결과에 대한 분석과 이해를 중심으로 학습 하시기 바랍니다.
- 제공된 2018_fall_session_2_task.ipynb 파일에 코드를 작성하고 설명을 단 뒤 ipynb 파일을 제출해주세요.
- 제출시간은 진행상황을 봐가며 정하도록 하겠습니다.
- 다음의 이메일로 choid@snu.ac.kr (최대영 조교) 로 제출하시고 보내실 때 반드시 조별 대표자 성명을 기록해주시기 바랍니다.